

Is L^2 Physics-Informed Loss Always Suitable for Training Physics-Informed Neural Network?

Chuwei Wang*
Shanda Li*
Di He
Liwei Wang



Background

Partial Differential Equation (PDE)

$$\begin{cases} \mathcal{L}u(x) = \varphi(x) & x \in \Omega \subset \mathbb{R}^n \\ \mathcal{B}u(x) = \psi(x) & x \in \partial\Omega, \end{cases}$$

\mathcal{L} : PDE operator

\mathcal{B} : boundary/initial conditions

x : spatiotemporal-dependent variable

Physics-Informed Neural Network (PINN)

Train an NN, u_θ , to represent the PDE solution by minimizing the Physics-Informed Loss:

$$\ell_{\Omega,p}(u) = \|\mathcal{L}u(x) - \varphi(x)\|_{L^p(\Omega)}^p,$$

$$\ell_{\partial\Omega,p}(u) = \|\mathcal{B}u(x) - \psi(x)\|_{L^p(\partial\Omega)}^p.$$

➤ $p = 2$ for vanilla PINN.

Probing PINN training

- Zero training loss \Leftrightarrow Learned solution is exactly accurate
- Small but non-zero loss \Rightarrow ? (This work)

Does a learned solution with a small Physics-Informed Loss always corresponds to a good approximator of the exact solution?

No!

Hamilton-Bellman-Jacobi (HJB) Equation

$$\begin{cases} \mathcal{L}_{\text{HJB}}u := \partial_t u(x, t) + \frac{1}{2}\sigma^2 \Delta u(x, t) - \sum_{i=1}^n A_i |\partial_{x_i} u|^{c_i} = \varphi(x, t) \\ \mathcal{B}_{\text{HJB}}u := u(x, T) = g(x) \end{cases}$$

- Most important PDE in stochastic control.
- Representative of high-dimensional nonlinear PDE.

Empirical finding: Relative error remains large in spite of the small training loss.

Iteration	1000	2000	3000	4000	5000
L^2 Loss	0.098	0.088	0.070	0.584	0.041
L^1 Relative Error	6.18%	5.36%	3.86%	3.94%	3.47%
$W^{1,1}$ Relative Error	17.53%	17.67%	14.83%	14.40%	11.31%

PINN through the lens of stability

Stability of PDE

[Definition] A PDE is (Z_1, Z_2, Z_3) -stable, if $\|u^*(x) - u(x)\|_{Z_3} = O(\|\mathcal{L}u(x) - \varphi(x)\|_{Z_1} + \|\mathcal{B}u(x) - \psi(x)\|_{Z_2})$ as $\|\mathcal{L}u(x) - \varphi(x)\|_{Z_1}, \|\mathcal{B}u(x) - \psi(x)\|_{Z_2} \rightarrow 0$, where Z_1, Z_2, Z_3 are three Banach spaces and u^* is the exact solution.

➤ Physics-Informed loss functions corresponding to $\|\cdot\|_{Z_1}$ and $\|\cdot\|_{Z_2}$ help to obtain u_θ which is **provably** close to the exact solution.

➤ PINN training with **L^2 Physics-Informed Loss** is suitable only when a PDE is **(L^2, L^2, Z) -stable** for some Banach space Z .

Stability property of HJB Equation

[Theorem 1] (Stability of HJB equations)
A large class of n -dimensional HJB Equation is $(L^p, L^q, W^{1,1})$ -stable if $p > n, q > kn$ (k depends on the equation).

[Theorem 2] (Instability of HJB equations)

There exists an instance of HJB Equation, whose exact solution is u^* , such that **for any** $\varepsilon > 0, A > 0, r \geq 1, m \in \mathbb{N}$ and $p \in [1, n/4]$, there exists a smooth function u which satisfies

- $\|\mathcal{L}_{\text{HJB}}u - \varphi\|_{L^p(\mathbb{R}^n \times [0, T])} < \varepsilon, \mathcal{B}_{\text{HJB}}u = \mathcal{B}_{\text{HJB}}u^*$, and $\text{supp}(u - u^*)$ is compact,
- $\|u - u^*\|_{W^{m,r}(\mathbb{R}^n \times [0, T])} > A$.

➤ We characterize the error using Sobolev norm as the gradient of the solution is also important for HJB Equation.

The distance between u_θ and u^ , ∇u_θ and ∇u^* can be arbitrarily large even though the L^2 loss is small!*

Minimizing L^∞ loss using adversarial training

New training objective for PINN

Theoretical results inspires us to minimize L^∞ Physics-Informed Loss

$$\ell_\infty(u) = \sup_{x \in \Omega} |\mathcal{L}u(x) - \varphi(x)| + \lambda \sup_{x \in \partial\Omega} |\mathcal{B}u(x) - \psi(x)|$$

Algorithm 1 L^∞ Training for Physics-Informed Neural Networks

Input: Target PDE (Eq. (1)); neural network u_θ ; initial model parameters θ

Output: Learned PDE solution u_θ

Hyper-parameters: Number of total training iterations M ; number of iterations and step size of inner loop K, η ; weight for combining the two loss term λ

```

1: for  $i = 1, \dots, M$  do
2:   Sample  $x^{(1)}, \dots, x^{(N_1)} \in \Omega$  and  $\tilde{x}^{(1)}, \dots, \tilde{x}^{(N_2)} \in \partial\Omega$ 
3:   for  $j = 1, \dots, K$  do
4:     for  $k = 1, \dots, N_1$  do
5:        $x^{(k)} \leftarrow \text{Project}_\Omega \left( x^{(k)} + \eta \text{sign} \nabla_x (\mathcal{L}u_\theta(x^{(k)}) - \varphi(x^{(k)}))^2 \right)$ 
6:     for  $k = 1, \dots, N_2$  do
7:        $\tilde{x}^{(k)} \leftarrow \text{Project}_{\partial\Omega} \left( \tilde{x}^{(k)} + \eta \text{sign} \nabla_x (\mathcal{B}u_\theta(\tilde{x}^{(k)}) - \psi(\tilde{x}^{(k)}))^2 \right)$ 
8:    $g \leftarrow \nabla_\theta \left( \frac{1}{N_1} \sum_{i=1}^{N_1} (\mathcal{L}u_\theta(x^{(i)}) - \varphi(x^{(i)}))^2 + \lambda \cdot \frac{1}{N_2} \sum_{i=1}^{N_2} (\mathcal{B}u_\theta(\tilde{x}^{(i)}) - \psi(\tilde{x}^{(i)}))^2 \right)$ 
9:    $\theta \leftarrow \text{Optimizer}(\theta, g)$ 
10: return  $u_\theta$ 

```

Inner loop:
Gradient ascend
for data points

Outer loop: Gradient descend for NN parameters

Experiments

High-dimensional LQG Problem

$$\begin{cases} \partial_t u(x, t) + \Delta u(x, t) - \mu \|\nabla_x u(x, t)\|^2 = 0 \\ u(x, T) = g(x) \end{cases}$$

Method	Relative error for $n = 100$		
	L^1	L^2	$W^{1,1}$
Original PINN [28]	3.47%	4.25%	11.31%
Adaptive time sampling [35]	3.05%	3.67%	13.63%
Learning rate annealing [34]	11.09%	11.82%	33.61%
Curriculum regularization [17]	3.40%	3.91%	9.53%
Adversarial training (ours)	0.27%	0.33%	2.22%

More HJB Equation variants

Method	$c = 1.25$	$c = 1.5$	$c = 1.75$
Original PINN [28]	1.11%	3.82%	2.73%
Adaptive time sampling [35]	1.18%	2.34%	7.94%
Learning rate annealing [34]	0.98%	1.13%	1.06%
Curriculum regularization [17]	6.27%	0.37%	3.51%
Adversarial training (ours)	0.61%	0.15%	0.29%

Solution visualization

